APPROXIMATING LIKELIHOOD RATIOS WITH CALIBRATED DISCRIMINATIVE CLASSIFIERS

By Kyle Cranmer[*]

*Center for Cosmology and Particle Physics[†]
and Center for Data Science[‡] New York University*

In particle physics likelihood ratio tests are established tools for statistical inference. These tests are complicated by the fact that computer simulators are used as a generative model for the data, but they do not provide a way to evaluate the likelihood function. We demonstrate how discriminative classifiers can be used to approximate the likelihood function when a generative model for the data is available for training and calibration. This offers an approach to parametric inference when simulators are used that is complementary to approximate Bayesian computation.

1. Introduction. The likelihood function is the central object that summarizes the information from an experiment needed for inference of model parameters. The likelihood function is key to Bayesian inference, and many areas of science that report the results of classical hypothesis tests or confidence intervals use the (generalized or profile) likelihood ratio as a test statistic. It is increasingly common that a simulator (or generative model) is used to describe complex processes that tie parameters θ of an underlying theory and measurement apparatus to high-dimensional observations x. Directly evaluating the likelihood ratio in these cases is often impossible or is computationally impractical. Approximate Bayesian Computation (ABC) is one approach to parameter inference in this simulation-based or likelihood-free setting [Rubin (1984); Tavaré et al. (1997); Marin et al. (2011)]. Here we consider an alternative approach that can also be used in a classical setting where a prior over the parameters is not available. In particular, we demonstrate how discriminative classifiers can be used to construct equivalent likelihood ratio tests when a generative model for the data is available for training and calibration.

As a concrete example, consider searches for new particles at the Large Hadron Collider. The simulator that is sampling from $p(x|\theta)$ is based on quantum field theory, a detailed simulation of the particle detector, and

---

data processing algorithms that transform raw sensor data into the feature vector x [Sjostrand, Mrenna and Skands (2006); Agostinelli et al. (2003)].

The ATLAS and CMS experiments have published hundreds papers where the final result was formulated as a hypothesis test or confidence interval using a generalized likelihood ratio test [Cowan et al. (2010)]. This includes the discovery of the Higgs boson [The ATLAS Collaboration (2012a); The CMS Collaboration (2012b)] and subsequent measurement of its properties.

The bulk of the likelihood ratio tests at the LHC are based on the distribution of a single event-level feature that discriminates between a hypothesized process of interest (labeled *signal*) and various other processes (labeled *background*). Typically, pseudo-data from the simulator are used to approximate the density at various parameter points, and an interpolation algorithm is used to approximate the parametrized model [Cranmer et al. (2012)].

To improve the statistical power of these tests, hundreds of these searches have utilized supervised learning to train discriminative classifiers that take advantage of a high dimensional feature vector x. Within High Energy Physics (HEP) libraries such as TMVA that implement conventional techniques like the multi-layer perceptron and boosted decision trees [Hocker et al. (2007)] are commonly used. Recently, there has been progress in using deep networks [Baldi, Sadowski and Whiteson (2014)] and a NIPS workshop synthesizing the lessons learned during HiggsML [Kégl et al.], the largest Kaggle challenge in history.

While classification accuracy can lead to optimal approaches for simple hypothesis tests [Dempster and Schatzoff (1965)], that is no longer true in the context of parameter estimation or composite hypothesis tests with nuisance parameters. As noted in [Whiteson and Whiteson (2007)], "such methods are suboptimal because they assume that the selector with the highest classification accuracy will yield a mass measurement with the smallest statistical uncertainty." The key distinction is that evaluating the loss for classification is composed of many per-event operations, while evaluating the loss for a mass measurement (e.g. the variance of an estimator for the mass parameter) is a per-experiment operation involving a data set with many events. They went on to demonstrate a computationally intensive stochastic optimization technique based on the per-experiment loss out performed the two stage selection-estimation process.

The initial motivation for this work was to extend the typical usage of discriminative classifiers in HEP to be robust to nuisance parameters in the simulators. The scope of the result expanded once it became clear that this offers a way to approximate the likelihood function $p(x|\theta)$ in what is typically considered the likelihood-free setting. This approach is complementary to

ABC as it does not require a prior over the parameters and can also be used in the classical (frequentist) setting. A strength of this approach is that it separates the quality of the approximation of the target likelihood from the quality of the calibration. In Section 7 we discuss the scheme sketched by Neal (2007) that also suggests using a classifier as a dimensionality reduction map to aid in the estimation of the likelihood function.

1.1. *Notation and Assumptions.* We use the following notation:

- x: a vector of features for an event
- D: a data set of $D = \{x_1, \ldots, x_n\}$, where x_e are assumed to be i.i.d.
- θ: parameters of a statistical model
- $p(x|\theta)$: probability density (simulation-based model) for x given θ
- y: a class label used for training a classifier.
- $s(x; \theta_0, \theta_1)$: real-valued discriminative classification score, parametrized by θ_0 and θ_1
- $p(s_{\theta_0, \theta_1}|\theta)$: The probability density for $s(x; \theta_0, \theta_1)$ implied by $p(x|\theta)$

We will assume the x_e are i.i.d., so that $p(D|\theta) = \prod_{e=1}^n p(x_e|\theta)$.

1.2. *Prelude.* In the setting where one is interested in simple hypothesis testing between a null $\theta = \theta_0$ against an alternate $\theta = \theta_1$, the Neyman-Pearson lemma states that the likelihood ratio

$$(1.1) \qquad T(D; \theta_0, \theta_1) = \prod_{e=1}^n \frac{p(x_e|\theta_0)}{p(x_e|\theta_1)}$$

is the most powerful test statistic. In order to evaluate $T(D)$, one must be able to evaluate the probability density $p(x|\theta)$ at any value x. However, it is increasingly common in science that one has a complex simulation that can act as generative model for $p(x|\theta)$, but one cannot evaluate the density directly. For instance, this is the case high energy physics where the simulation of particle detectors can only be done in the 'forward mode'. This same setting has been considered by Scott and Nowak (2005), Xin Tong (2013), and Neal (2007).

The main result of this paper is to generalize the observation that one can form an equivalent test based on

$$(1.2) \qquad T'(D; \theta_0, \theta_1) = \prod_{e=1}^n \frac{p(s_e|\theta_0)}{p(s_e|\theta_1)}$$

if

$$(1.3) \qquad s_e = s(x_e; \theta_0, \theta_1) = m\left(p(x_e|\theta_0)/p(x_e|\theta_1)\right)$$

where m is any strictly increasing or decreasing function. This result will be proven below. This allows us to recast the original likelihood ratio test into an alternate form in which supervised learning is used to train the discriminative classifier $s(x; \theta_0, \theta_1)$. The discriminative classifier can be trained with data $(x, y = 0)$ generated from $p(x|\theta_0)$ and $(x, y = 1)$ generated from $p(x|\theta_1)$. In Section 4 we extend this result to generalized likelihood ratio tests, where it will be useful to have the classifier parametrized in terms of (θ_0, θ_1).

Here we see that the original goal for simple hypothesis testing (i.e. to make a decision to accept or reject the null hypothesis based on the entire data set D) has been reformulated into a per-event classification problem. This follows from the fact that we assume the x_e to be i.i.d.

1.3. *Comments on classification and frequentist hypothesis tests.* Vast literature exists around generative and discriminative classifiers [Ng and Jordan (2002)]. Typically, generative classifiers learn a model for the joint probability $p(x, y)$, of the inputs x and the classification label y, and predict $p(y|x)$ via Bayes rule. In contrast, discriminative classifiers model the posterior $p(y|x)$ directly. For classification tasks, one then thresholds on $p(y|x)$. In both cases this description in terms of a posterior requires a prior distribution for $p(y)$, which is either modeled explicitly or learned from the training data. This familiar formulation of classification may lead to some confusion in the setting of the current work.

The first possible source of confusion we wish to avoid is that here $p(x|\theta)$ is a *generative statistical model* for the features x, not a generative classifier. We think of the $p(x|\theta)$ along the lines of a traditional scientific theory or simulator, able to make predictions about x and being motivated by domain-specific considerations.

The second possible source of confusion is that we are not directly interested in calibrating the classification score in terms of a per-event posterior probability $p(y|x)$. Instead, we are interested in the approximation of the per-experiment likelihood function or likelihood ratio, which might be used for several purposes, including the calculation of p-values.

Lastly, in the setting of frequentist hypothesis tests and confidence intervals, we do not have a prior $\pi(\theta)$. While we can use the generative models to produce training data $(x, y = 0)$ generated from $p(x|\theta_0)$ and $(x, y = 1)$ generated from $p(x|\theta_1)$, the relative mix $p(y)$ is arbitrary. When $p(y = 0) = p(y = 1) = 1/2$, then

$$(1.4) \qquad p(y = 1|x) = \frac{p(x|y = 1)}{p(x|y = 0) + p(x|y = 1)} = \frac{p(x|\theta_1)}{p(x|\theta_0) + p(x|\theta_1))} \,,$$

which is monotonic with the desired likelihood ratio $p(x|\theta_1)/p(x|\theta_0)$. Since the prior $p(y)$ is not needed for the target likelihood ratio test and because the classifier score $p(y|x)$ may not be well calibrated, we choose to denote the classifier score $s(x)$ and simply think of it as a deterministic dimensionality reduction map $s : X \to \mathbb{R}$. Similar points have been made by Scott and Nowak (2005) and Neal (2007).

2. Dimensionality reduction and calibration. We are interested in reformulating the target likelihood ratio

$$(2.1) \qquad \ln T(D; \theta_0, \theta_1) = \sum_{e=1}^{n} \underbrace{\log \left[\frac{p(x_e|\theta_0)}{p(x_e|\theta_1)} \right]}_{q(x_e)} .$$

Here we see that the test statistic T for the experiment is composed of a sum over events of the per-event function $q(x)$. A sum over a monotonic, but non-linear function of $q(x)$ would not lead to an equivalent statistic.

The important part of the per-event function $q(x)$ is that it defines iso-contours in the feature space x. As we will show, our goal is to learn a monotonic function of $p(x|\theta_0)/p(x|\theta_1)$, which will share the same iso-contours. Then the remaining challenge is to find the appropriate monotonic function that gives back a linear function of $q(x)$. Our claim is that the generative model $p(x|\theta)$ can be used to calibrate the density $p(s|\theta)$ and that

$$(2.2) \qquad \ln T'(D; \theta_0, \theta_1) = \sum_{e=1}^{n} \underbrace{\log \left[\frac{p(s_e|\theta_0)}{p(s_e|\theta_1)} \right]}_{q(s_e)} ,$$

leads to an equivalent statistic.

For notational simplicity, let $p_0(x) = p(x|\theta_0)$, $p_1(x) = p(x|\theta_1)$, and $s(x) = s(x; \theta_0, \theta_1)$. The distribution of x totally determines the distribution of s. In the application at hand, the function s maps a high-dimensional feature vector x to $\mathbb{R}^+$. Let Ω_{s^*} be the level set $\{x \mid s(x; \theta_0, \theta_1) = s^*\}$ and $\hat{n} = \nabla s(x)/|\nabla s(x)|$ be the orthonormal vector to Ω_{s^*} at the point x.

We need to show that for all x, the density

$$(2.3) \qquad p(q_x|\theta) = \int dx \delta(q_x - q_x(x)) p(x|\theta)/|\hat{n} \cdot \nabla q_x|$$

is equal to the density

$$(2.4) \qquad p(q_s|\theta) = \int dx \delta(q_s - q_s(s(x))) \, p(x|\theta) \, /|\hat{n} \cdot \nabla q_s| .$$

It is sufficient to show that $q_x(x) = q_s(s(x))$. The function $q_s(s)$ is based on the induced densities $p_0(s)$ and $p_1(s)$. The induced density $p_1(s)$ is given by

$$(2.5) \qquad p_1(s^*) = \int dx \delta(s^* - s(x)) p_1(x) = \int d\Omega_{s^*} p_1(x)/|\hat{n} \cdot \nabla s|$$

and a similar equation for $p_0(s)$.

Theorem 1: We have the following equality

$$(2.6) \qquad\qquad \frac{p_1(s(x))}{p_0(s(x))} = \frac{p_1(x)}{p_0(x)} \ .$$

Proof For $x \in \Omega_{s^*}$, we can factor out of the integral the constant $p_1(x)/p_0(x)$. Thus

$$(2.7) \qquad p_1(s^*) = \int d\Omega_{s^*} p_1(x)/|\hat{n} \cdot \nabla s| = \frac{p_1(x)}{p_0(x)} \int d\Omega_{s^*} p_0(x)/|\hat{n} \cdot \nabla s| \ ,$$

and the integrals cancel in the likelihood ratio

$$(2.8) \qquad \frac{p_1(s^*)}{p_0(s^*)} = \frac{p_1(x)}{p_0(x)} \frac{\int d\Omega_{s^*} p_0(x)/|\hat{n} \cdot \nabla s|}{\int d\Omega_{s^*} p_0(x)/|\hat{n} \cdot \nabla s|} = \frac{p_1(x)}{p_0(x)} \qquad\qquad \forall x \in \Omega_{s^*}.$$

One can think of the ratio $p_1(s)/p_0(s)$ as a way of calibrating the the discriminative classifier and correcting for the monotonic transformation m of the desired likelihood ratio as in Eq. 1.3.

3. Embedding the classifier into the likelihood. Thus far we have shown that the target likelihood ratio $p(x|\theta_0)/p(x|\theta_1)$ with high dimensional features x can be reproduced via the univariate densities $p(s|\theta_0)/p(s|\theta_1)$ if the classifier $s(x|\theta_0, \theta_1)$ is a strictly increasing function of $p(x|\theta_0)/p(x|\theta_1)$. We now generalize from the ratio of two simple hypotheses specified by θ_0 and θ_1 to the case where θ are continuous model parameters. We postpone the practicalities of training the classifier and estimating the density to Section 5 and continue in the likelihood-free setting with idealized classifiers and their densities.

In the case of a fixed classifier $s(x)$ it is possible to compute $s_e = s(x_e)$ for the observed data and never refer back to the original features x_e. In the parametrized setting it is not possible to pre-compute $s(x_e; \theta_0, \theta_1)$ since θ_0 and θ_1 are unknown.

The critical observation is that if we postpone the evaluation of the classifier to the stage of evaluating the enveloping likelihood ratio, then we can

identify the value of the parameters that are being compared in the likelihood ratio with the values used as input to $s(x; \theta_0, \theta_1)$.

$$(3.1) \qquad T(D; \theta_0, \theta_1) = \prod_e \frac{p(x_e|\theta_0)}{p(x_e|\theta_1)} = \prod_e \frac{p(s(x_e; \theta_0, \theta_1)|\theta_0)}{p(s(x_e; \theta_0, \theta_1)|\theta_1)} .$$

This is equivalent to approximating the likelihood function for θ_0 when θ_1 is held fixed.

4. Composite hypotheses and the generalized likelihood ratio. In the case of composite hypotheses $\theta \in \Theta_0$ against an alternative $\theta \in \Theta_0^C$, the generalized likelihood ratio[1] test is commonly used

$$(4.1) \qquad \lambda(\Theta_0) = \frac{\sup_{\theta \in \Theta_0} p(D|\theta)}{\sup_{\theta \in \Theta} p(D|\theta)} .$$

This generalized likelihood ratio can be used both for hypothesis tests in the presence of nuisance parameters or to create confidence intervals with or without nuisance parameters. Often, the parameter vector is broken into two components $\theta = (\mu, \nu)$, where the μ components are considered parameters of interest while the ν components are considered nuisance parameters. In that case Θ_0 corresponds to all values of ν with μ fixed.

Denote the maximum likelihood estimator

$$(4.2) \qquad \hat{\theta} = \arg\max_\theta p(D|\theta)$$

and the conditional maximum likelihood estimator

$$(4.3) \qquad \hat{\hat{\theta}} = \arg\max_{\theta \in \Theta_0} p(D|\theta) .$$

It is not obvious that if we are working with the distributions $p(s|\theta)$ (for some particular $s(x; \theta_0, \theta_1)$ comparison) that we can find the same estimators. Fortunately, there is a construction based on $p(s|\theta)$ that works. The maximum likelihood estimate of Eq. 4.2 is the same as the value that maximizes the likelihood ratio with respect to $p(D|\theta_1)$ for some fixed value of θ_1. This allows us to use Theorem 1 to reformulate the maximum likelihood estimate

$$(4.4) \qquad \hat{\theta} = \arg\max_\theta \sum \ln \frac{p(x_e|\theta)}{p(x_e|\theta_1)} = \arg\max_\theta \sum \ln \frac{p(s(x_e; \theta, \theta_1)|\theta)}{p(s(x_e; \theta, \theta_1)|\theta_1)} .$$

It is important that we include the denominator $p(s(x_e; \theta, \theta_1)|\theta_1)$ because this cancels Jacobian factors that vary with θ.

---

[1]Also known as the profile likelihood ratio.

5. Learning the correct mapping and its distribution.   In order for this approach to be useful in the likelihood-free setting where it is not possible to evaluate the density $p(x|\theta)$, we need to be able to approximate both $s(x|\theta_0, \theta_1)$ and $p(s|\theta)$ based on samples $\{x_i\}$ drawn from the generative model $p(x|\theta)$.

Denote the approximate dimensionality reduction map $\hat{s}(x; \theta_0, \theta_1)$ and its distribution $\hat{p}(\hat{s}|\theta)$. In general we will be interested in the machine learning problem that approximates these distributions based on samples $\{x_i\}$ drawn from the generative model $p(x|\theta)$. The first step in this direction is to confirm that a discriminative classifier obtained from common supervised learning algorithms will yield a function that is one-to-one with $p(x|\theta_0)/p(x|\theta_1)$.

5.1. *The fixed discriminative classification setting.*   For fixed θ_0 and θ_1 we can generate large samples from each model and train a classifier. To be concrete, let us use $p(x|\theta_0)$ to generate training data $(x_i, y_i = 0)$ and $p(x|\theta_1)$ to generate training data $(x_i, y_i = 1)$. With balanced training data $(p(y = 1) = p(y = 0) = 1/2)$ a quadratic loss function will lead to classifiers that approximate the regression function $\hat{s}(x) \approx p(y|x) = p(x|\theta_1)/(p(x|\theta_0) + p(x|\theta_1))$, which is monotonic with the desired per-event likelihood ratio $q(x)$. Thus, standard supervised learning algorithms with various surrogate loss functions lead to discriminative classifiers that approximate a monotonic function of per-event likelihood ratio $q(x)$.

Once the classifier is trained, we can use the generative model together with a univariate density estimation technique (e.g. histograms or kernel density estimation) to approximate $\hat{p}(\hat{s}|\theta)$ for specific parameter points. In order to construct an approximate likelihood function continuous in the parameters θ, HEP experiments have typically used some form of interpolation of the univariate score density over the parameters θ [Cranmer et al. (2012)].

5.2. *Training a parametrized classifier.*   We are left with the practical question of how to train a family of discriminative classifiers parametrized by θ_0 and θ_1, the parameters associated to the null and alternate hypotheses, respectively. While this could be done independently for all θ_0 and θ_1, it is desirable and convenient to have a smooth evolution of the classification score as a function of the parameters. Thus, we anticipate a single learning stage based on training data with input $(x, \theta_0, \theta_1)_i$ and target y_i. Somewhat unusually, the unknown values of the parameters are taken as input to the classifier, and are latent variables whose values will be specified via the enveloping (generalized) likelihood ratio of Eq. 2.2. We denote the learned family of classifiers $\hat{s}(x; \theta_0, \theta_0)$, and anticipate the training based roughly on the following algorithmic flow.

Algorithm 1 Training of the parametrized classifier.

---

initialize trainingData to {}
for θ_0 in Θ **do**
   for θ_1 in Θ **do**
      generate $x_i^0 \sim p(x|\theta_0)$
      append $\{(x_i^0, \theta_0, \theta_1, y = 0)\}$ to trainingData
      generate $x_i^1 \sim p(x|\theta_1)$
      append $\{(x_i^1, \theta_0, \theta_1, y = 1)\}$ to trainingData
   end for
end for
use trainingData to learn $\hat{s}(x; \theta_0, \theta_1)$

While the function $p(x|\theta_1)/(p(x|\theta_0) + p(x|\theta_1))$ will minimize the expected squared loss based on training data produced according to Algorithm 1, it is not clear how training data from $\theta_0' \neq \theta_0$ and $\theta_1' \neq \theta_1$ will influence a real world classifier with finite capacity. This is left as an area for future work.

5.3. *Practicalities of embedding the classifier into the likelihood.*   In both the fixed and parametrized cases, constructing the density $\hat{p}(\hat{s}|\theta)$ requires running the generative model at θ. In the context of a likelihood fit this would mean that the optimization algorithm that is trying to maximize the likelihood with respect to θ needs access to the generative model $p(x|\theta)$. This can be impractical when the generative model is computationally expensive or has high-latency (for instance some human intervention is required to reconfigure the generative model). In HEP with a fixed classifier, it has become common to interpolate the distribution between discrete values of θ in order to produce a continuous parametrization for $\hat{p}(\hat{s}|\theta)$.

In the case of a fixed classifier $\hat{s}(x)$ it is possible to pre-compute $\hat{s}_e = \hat{s}(x_e)$ and never refer back to the original features x_e. In the parametrized setting it is not possible to pre-compute $\hat{s}(x_e; \theta_0, \theta_1)$ since one does not know the true values of θ_0 and θ_1; however, we can embed the classifier as in Eq. 3.1.

A concrete realization of this has been performed for probability models implemented with the `RooFit` probabilistic programing language and classifiers implemented with `scikit-learn` and `TMVA` [Verkerke and Kirkby (2003); Pedregosa et al. (2011); Hocker et al. (2007)].

One can easily imagine a number of approaches to embedding the classifier and estimating the density $\hat{p}(\hat{s}|\theta)$ and the relative merits of those approaches will depend critically on the dimensionality of θ and the computational cost of the generative model. We leave a more general strategy for this overarching optimization problem as an area of future work.

6. Applications in HEP. In HEP we are often searching for some class of events, generically referred to as *signal*, in the presence of a separate class of *background* events. For each event we measure some quantities x that have corresponding distributions $p_b(x|\nu)$ for background and $p_s(x|\nu)$ for signal, where ν are nuisance parameters describing uncertainties in the underlying physics prediction or response of the measurement device. The total model is a mixture of the signal and background, and μ is the mixture coefficient associate to the signal component.

$$(6.1) \qquad p(D \,|\, \mu, \nu) = \prod_{e=1}^{n} \left[\mu p_s(x_e \,|\, \nu) + (1 - \mu) \, p_b(x_e \,|\, \nu) \right] \,,$$

New particle searches at the LHC are typically framed as hypothesis test where the null corresponds to $\mu = 0$, and the generalized likelihood ratio (profiling over ν as in Eq. 4.3) is used as a test statistic [Cowan et al. (2010)]. This approach was used for the discovery of the Higgs boson [The ATLAS Collaboration (2012a); The CMS Collaboration (2012b)].

6.1. *Typical usage pattern.* In this setting, large samples of pseudo-data $\{x_i, y_i\}$ generated with some nominal values of the parameters ν_0, where $y = 0$ corresponds to the background density $p_b(x|\nu_0)$ and $y = 1$ corresponds to signal density $p_s(x|\nu_0)$. Importantly, the $y = 1$ label corresponds to the signal only, and not to the alternate signal-plus-background hypothesis. The resulting classifier approximates the regression function $p_s(x|\nu_0)/(p_s(x|\nu_0) + p_b(x|\nu_0))$, which is one to one with the likelihood ratio of the null to the alternate $p(x|\mu = 0, \nu_0)/p(x|\mu, \nu_0)$ for all μ. Associating the $y = 1$ label to the signal component has advantages because it helps the classifier focus its capacity on the relevant regions in the feature space, particularly when the signal is a very small perturbation to the background.

Once the classifier is trained, large samples of pseudo-data are drawn from $p_s(x|\nu)$ and $p_b(x|\nu)$ and we estimate the distributions $\hat{p}_s(\hat{s}|\nu)$ and $\hat{p}_b(\hat{s}|\nu)$ continuously parametrized in ν. An example of the distributions of $\hat{s}$ for the signal and background events with $\nu = \nu_0$ is shown in Figure 1.

These steps feed into a subsequent statistical test based on the observed data $D = (x_1, \ldots, x_n)$. For each event, the classifier is evaluated and one performs inference on a parameter μ related to the presence of the signal

FIG 1. *Left: an example of the distributions $p_b(\hat{s}|\nu)$ and $p_s(\hat{s}|\nu)$ when the classifier s is a boosted-decision tree (BDT). Right: the corresponding ROC curve (right) for this and other classifiers. (Figures taken from TMVA manual.)*

contribution. In particular, one forms the statistical model[2]

$$(6.2) \qquad p(D \,|\, \mu, \nu) = \prod_{e=1}^{n} \left[\mu \hat{p}_s(\hat{s}(x_e) \,|\, \nu) + (1 - \mu) \, \hat{p}_b(\hat{s}(x_e) \,|\, \nu) \right] .$$

6.2. *Comments on typical usage of machine learning in HEP.* Nuisance parameters are an after thought in the typical usage of machine learning in HEP. In fact, most discussions would related to the training and optimizing the classifier only consider $p_b(x)$ and $p_s(x)$ with $\nu = \nu_0$ being implicit. However, as experimentalists we know that we must account for various forms of systematic uncertainty, parametrized by nuisance parameters ν. In practice, we take the classifier as fixed and then propagate uncertainty through the classifier as in Eq. 6.2. Building the distribution $p(\hat{s}|\nu)$ for values of ν other than the nominal ν_0 used to train the classifier can be thought of as a calibration necessary for classical statistical inference; however, this classifier is clearly not optimal for $\nu \neq \nu_0$.

6.3. *A more powerful approach.* The standard use of machine learning in HEP can be improved by training a parametrized, discriminative classifier corresponding to the generalized likelihood ratio test

$$(6.3) \qquad \lambda(\mu) = \frac{p(D|\mu, \hat{\hat{\nu}})}{p(D|\hat{\mu}, \hat{\nu})} ,$$

---

[2]Sometimes there is an additional Poisson term when expected number of signal and background events is known, which is referred to as an extended likelihood or marked Poisson model.

following the approach outlined in Section 4.

There is an interesting distinction between this approach and the standard use in which the classifier is trained for a fixed ν_0. In the standard use one trains a classifier for signal vs. background, which is equivalent (in an ideal setting) to training a classifier for null (background-only) vs. alternate (signal-plus-backgound) since the resulting regression functions are one-to-one with each other. In contrast, in the case of the generalized likelihood ratio test

$$(6.4) \qquad \frac{p(x|0, \hat{\nu})}{p(x|\hat{\mu}, \hat{\nu})} = \frac{p_b(x|\hat{\nu})}{\hat{\mu} p_s(x_e \,|\, \hat{\nu}) + (1 - \hat{\mu}) \, p_b(x_e \,|\, \hat{\nu})} \;,$$

the background components don't cancel and there is an additional term $p_b(x|\hat{\nu})/p_b(x|\hat{\nu})$. In practice, with classifiers of finite capacity, there will be some tradeoff between taking into account this additional term and the more challenging learning problem when μ is very small.

6.4. *Decomposing tests between mixture models into their components.* In this section we generalize the capacity focusing technique of training classifiers to discriminate between components of a mixture model. First, we generalize Eq. 6.1 to a mixture model of several components

$$(6.5) \qquad p(x|\theta) = \sum_c w_c(\theta) p_c(x|\theta) \;.$$

It is possible to re-write the target likelihood ratio between two mixture models in terms of pairwise classification problems.

$$(6.6) \qquad \frac{p(x|\theta_0)}{p(x|\theta_1)} \;=\; \frac{\sum_c w_c(\theta_0) p_c(x|\theta_0)}{\sum_{c'} w_{c'}(\theta_1) p_{c'}(x|\theta_1)}$$

$$(6.7) \qquad \qquad =\; \sum_c \left[\sum_{c'} \frac{w_{c'}(\theta_1)}{w_c(\theta_0)} \frac{p_{c'}(x|\theta_1)}{p_c(x|\theta_0)} \right]^{-1}$$

$$(6.8) \qquad \qquad =\; \sum_c \left[\sum_{c'} \frac{w_{c'}(\theta_1)}{w_c(\theta_0)} \frac{p_{c'}(s_{c,c',\theta_0,\theta_1}|\theta_1)}{p_c(s_{c,c',\theta_0,\theta_1}|\theta_0)} \right]^{-1}$$

The second line is a trivial, but useful decomposition into pair-wise classification between $p_{c'}(x|\theta_1)$ and $p_c(x|\theta_0)$. The third line uses Theorem 1 to relate the high-dimensional likelihood ratio into an equivalent calibrated likelihood ratio based on the univariate density of the corresponding classifier, denoted $s_{c,c',\theta_0,\theta_1}$. In the situation where the only free parameters of the model are the mixture coefficients w_c, then the distributions $p_c(s_{c,c',\theta_0,\theta_1}|\theta)$

are independent of θ and can be pre-computed (after training the discriminative classifier, but before evaluating the likelihood ratio). Equation 6.8 allows one to take advantage of both the parametrized classifier as in Eq. 6.4 and the capacity focusing technique in the typical HEP usage pattern.

6.5. *Measuring particle properties.*   While the original motivation for this work was to improve the treatment of systematic uncertainties in new particle searches by parametrizing the classifier in terms of the nuisance parameters ν, the same approach can be used to for parameter inference. In the case of new particle searches the parameter of interest is the mixture coefficient for the signal component $p_s(x|\nu)$. When measuring particle properties the distribution of the features also depend on parameters such as a particle's mass and quantum numbers. This is easily accommodated by extending $p_s(x|\nu) \rightarrow p_s(x|\theta)$, where θ includes both parameters of interest and nuisance parameters.

This formalism represents is a significant step forward in the usage of machine learning in HEP, where classifiers have always been used between two static classes of events and not parametrized explicitly in terms of the physical quantities we wish to measure. The work of [Whiteson and Whiteson (2007)] is the closest example as the stochastic optimization was directly trying to minimize the measurement uncertainty of a particle's mass; however, the resulting classifier was fixed. This approach also offers the advantage that it explicitly reformulates the per-experiment optimization to the per-event optimization, which is less computationally intensive.

Another approach that is similar in spirit is the so-called matrix element method, in which one directly computes an approximate likelihood ratio by performing a computationally intensive integral associated to the detector response [Volobouev (2011)]. In the approach considered in this paper, the detector response is naturally handled by the Monte Carlo sampling used in the simulation of the detector; however, that integral is intractable for the matrix element method. Even with drastic simplifications of the detector response, the matrix element method can take several minutes of CPU time per event to calculate the likelihood ratio $q(x)$. The work here can be seen as aiming at the same conceptual target, but utilizing machine learning to overcome the complexity of the detector simulation. It also offers enormous speed increase for evaluating the likelihood at the cost of an initial training stage. In practice, the matrix element method has only been used for searches and measurement of a single physical parameter (sometimes with a single nuisance parameter as in [Aaltonen et al. (2010)]).

Contemporary examples where the technique presented here could have

major impact on HEP include the measurement of coefficients to quantum mechanical operators describing the decay of the Higgs boson [Chen et al. (2015)] and, if we are so lucky, measurement of the mass of supersymmetric particles in cascade decays [Allanach et al. (2000)]. Both of these examples involve data sets with many events, each with a feature vector x that has on the order of 10 components, and a parameter vector θ with 5-10 parameters of interest and possibly many more nuisance parameters. The state of the art for the operator coefficients of the Higgs decay uses the so-called matrix element likelihood analysis (MELA) in which the equivalent of $s(x; \theta_0, \theta_1)$ is approximated by neglecting detector effects [Bolognesi et al. (2012)].

7. Related work. Scott and Nowak (2005) and Xin Tong (2013) consider the machine learning problem associated to Neyman-Pearson hypothesis testing. As in this work, they consider the situation where one does not have access to the underlying distributions, but only has i.i.d. samples from each hypothesis. This work generalizes that goal from the Neyman-Pearson setting to generalized likelihood ratio tests and emphasizes the connection with classification. Perhaps a formal treatment similar to the Neyman-Pearson case can be brought to bear in this more general setting. Jaakkola and Haussler (1998) explore a way of leveraging generative models to derive kernel functions for use in discriminative methods. This interesting work is distinct from the point made here in which the generative model is being used for the purpose of providing training data and calibration. Zadrozny and Elkan (2001) emphasize the importance of calibrated probability estimates from decision trees and naive Bayesian classifiers and investigate various approaches to achieve this. In contrast to that work, we are not interested in calibrated probability estimates for $p(y|x)$ for individual events, but instead we use the calibration to correct for non-linear transformations of the target likelihood ratio and, perhaps, to provide calibrated p-values based on those likelihood ratio tests. Ihler, Fisher and Willsky (2004) take on a different problem (tests of statistical independence) by using machine learning algorithms to find scalar maps from the high-dimensional feature space that achieve the desired statistical goal when the fundamental high-dimensional test is intractable.

Neal (2007) also considered the problem of approximating the likelihood function when only a generative model is available. That work sketches a scheme in which one uses a classifier with both x and θ as an input to serve as a dimensionality reduction map. The key distinction comes in the handling of θ. Neal says "we cannot use the classifier on real data, since we

don't know the correct value for $[\theta]$" and goes on to outline an approach where one uses regression on a per-event basis to estimate $\hat{\theta}(x)$ and perform the composition $s(x; \hat{\theta}(x))$ (much like profiling).[3] This can lead to a significant loss of information since (at least in most particle physics examples) a single event carries little information about the true value of θ, though the full data set D may be informative – for instance, a single observation would not be sufficient to estimate the variance of a distribution, though repeated observations would. The work of Neal correctly identifies this as an approximation of the target likelihood even in the case of a ideal classifier. In contrast, the approach described here does not eliminate the dependence of the classifier on θ.[4] Instead, we embed a parametrized classifier into the likelihood and postpone the evaluation of the classifier to the point of evaluation of the likelihood when θ is explicitly being tested. This avoids the loss of information that occurs from the regression step $\hat{\theta}(x)$ proposed by Neal and leads to Theorem 1, which is an exact result in the case of an ideal classifier. In both cases, the quality of the classifier is factorized from the calibration of its density.

8. Conclusions. We have outlined a technique to reformulate generalized likelihood ratio test over a high-dimensional data set with multiple events in terms of a univariate density of a classifier score. We have shown that a parametrized family of discriminative classifiers $s(x; \theta_0, \theta_1)$ trained and calibrated with a simulator $p(x|\theta)$ can be used to approximate the likelihood ratio $\prod_i p(x_i|\theta_0)/p(x_i|\theta_1)$ even when it is not possible to directly evaluate the likelihood $p(x|\theta)$. It offers an alternative to approximate Bayesian computation for parameter inference in the likelihood-free setting that can also be used in the frequentist formalism without specifying a prior over the parameters. A strength of this approach is that it separates the quality of the approximation of the target likelihood from the quality of the calibration. The former is related to the ability of supervised learning approaches to classification, which will continue to improve. The calibration procedure for a particular parameter point is fairly straight forward since it involves estimating a univariate density using a generative model of the data. The

---

[3]Neal considers a lower-dimensional 'bottlenecks' $\theta^* = g(\theta)$, which are not essential to the discussion here.

[4]As a technical point, in Neal's work, the focus is on approximating the likelihood function (up to a multiplicative constant), which is equivalent to evaluating the ratio with respect to a fixed θ_1 as in Eq. 4.4. In Neal's case, the dependence on θ is eliminated via $s(x; \hat{\theta}(x))$ and the map is constant; however, in this approach the map ratio is explicitly parametrized in terms of θ, so the ratio is important for canceling the corresponding Jacobian factors.

difficulty of the calibration stage is performing this calibration continuously in θ. Different strategies to this calibration are anticipated depending on the dimensionality of θ, the complexity of the resulting likelihood function, and the difficulties associated to running the simulator.

References.

AALTONEN, T. et al. (2010). Top Quark Mass Measurement in the Lepton + Jets Channel Using a Matrix Element Method and *in situ* Jet Energy Calibration. *Phys.Rev.Lett.* **105** 252001.

AGOSTINELLI, S. et al. (2003). GEANT4: A Simulation toolkit. *Nucl.Instrum.Meth.* **A506** 250-303.

ALLANACH, B. C., LESTER, C. G., PARKER, M. A. and WEBBER, B. R. (2000). Measuring sparticle masses in nonuniversal string inspired models at the LHC. *JHEP* **0009** 004.

BALDI, P., SADOWSKI, P. and WHITESON, D. (2014). Searching for Exotic Particles in High-Energy Physics with Deep Learning. *Nature Commun.* **5** 4308.

BOLOGNESI, S., GAO, Y., GRITSAN, A. V., MELNIKOV, K., SCHULZE, M. et al. (2012). On the spin and parity of a single-produced resonance at the LHC. *Phys.Rev.* **D86** 095031.

CHEN, Y., DI MARCO, E., LYKKEN, J., SPIROPULU, M., VEGA-MORALES, R. et al. (2015). 8D likelihood effective Higgs couplings extraction framework in $h \to 4\ell$. *JHEP* **1501** 125.

THE ATLAS COLLABORATION (2012a). Observation of a new particle in the search for the Standard Model Higgs boson with the ATLAS detector at the LHC. *Phys.Lett.* **B716** 1–29.

THE CMS COLLABORATION (2012b). Observation of a new boson at a mass of 125 GeV with the CMS experiment at the LHC. *Phys.Lett.* **B716** 30–61.

COWAN, G., CRANMER, K., GROSS, E. and VITELLS, O. (2010). Asymptotic formulae for likelihood-based tests of new physics. *Eur.Phys.J.* **C71** 1554.

CRANMER, K., LEWIS, G., MONETA, L., SHIBATA, A. and VERKERKE, W. (2012). HistFactory: A tool for creating statistical models for use with RooFit and RooStats. CERN-OPEN-2012-016, http://inspirehep.net/record/1236448.

DEMPSTER, A. P. and SCHATZOFF, M. (1965). Expected Significance Level as a Sensitivity Index for Test Statistics. *Journal of the American Statistical Association* **60** pp. 420-436.

HOCKER, A., STELZER, J., TEGENFELDT, F., VOSS, H., VOSS, K. et al. (2007). TMVA - Toolkit for Multivariate Data Analysis. *PoS* **ACAT** 040.

IHLER, A. T., FISHER, J. W. and WILLSKY, A. S. (2004). Nonparametric Hypothesis Tests for Statistical Dependency. *IEEE Transactions on Signal Processing* **52** 2234–2249. http://ieeexplore.ieee.org/lpdocs/epic03/wrapper.htm?arnumber=1315943.

JAAKKOLA, T. and HAUSSLER, D. (1998). Exploiting Generative Models in Discriminative Classifiers. http://citeseer.ist.psu.edu/viewdoc/summary?doi=10.1.1.44.7709.

KÉGL, B., ROUSSEAU, D., COWAN, G., GERMAIN, C. and GUYON, I. HepML Workshop https://sites.google.com/site/hepml14/home.

MARIN, J.-M., PUDLO, P., ROBERT, C. P. and RYDER, R. J. (2011). Approximate Bayesian computational methods. *Statistics and Computing* **22** 1167–1180.

NEAL, R. M. (2007). Computing likelihood functions for high-energy physics experiments when distributions are defined by simulators with nuisance parameters. In *Proceedings of PhyStat2007, CERN-2008-001* 111-118. http://inspirehep.net/record/776337.

NG, A. Y. and JORDAN, M. I. (2002). On Discriminative vs. Generative Classifiers: A comparison of logistic regression and naive Bayes. In *Advances in Neural Information Processing Systems 14* (T. G. DIETTERICH, S. BECKER and Z. GHAHRAMANI, eds.) 841–848. MIT Press.

PEDREGOSA, F., VAROQUAUX, G., GRAMFORT, A., MICHEL, V., THIRION, B., GRISEL, O., BLONDEL, M., PRETTENHOFER, P., WEISS, R., DUBOURG, V., VANDERPLAS, J., PAS-

sos, A., Cournapeau, D., Brucher, M., Perrot, M. and Duchesnay, E. (2011).
Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research* **12**
2825–2830.

Rubin, D. B. (1984). Bayesianly Justifiable and Relevant Frequency Calculations for the
Applied Statistician. *The Annals of Statistics* **12** 1151–1172.

Scott, C. and Nowak, R. (2005). A Neyman-Pearson approach to statistical learn-
ing. *IEEE Trans. Inform. Theory* **51** 3806–3819. http://citeseerx.ist.psu.edu/
viewdoc/summary?doi=10.1.1.103.6850&rank=5.

Sjostrand, T., Mrenna, S. and Skands, P. Z. (2006). PYTHIA 6.4 Physics and Man-
ual. *JHEP* **0605** 026.

Tavaré, S., Balding, D. J., Griffiths, R. C. and Donnelly, P. (1997). Inferring
coalescence times from DNA sequence data. *Genetics* **145** 505–18.

Xin Tong (2013). A Plug-in Approach to Neyman-Pearson Classification. *Journal of
Machine Learning Research* **14** 3011–3040.

Verkerke, W. and Kirkby, D. P. (2003). The RooFit toolkit for data modeling. *eConf*
C0303241 MOLT007.

Volobouev, I. (2011). Matrix Element Method in HEP: Transfer Functions, Efficiencies,
and Likelihood Normalization. http://arxiv.org/abs/1101.2259.

Whiteson, S. and Whiteson, D. (2007). Stochastic optimization for collision selection
in high energy physics. *Association for the Advancement of Artificial Intelligence* **292**.

Zadrozny, B. and Elkan, C. (2001). Obtaining calibrated probability estimates from
decision trees and naive Bayesian classifiers. In *In Proceedings of the Eighteenth
International Conference on Machine Learning* 609–616. Morgan Kaufmann http:
//citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.29.3039.

Center for Cosmology and Particle Physics
New York University
New York, NY 10003, USA   kyle.cranmer@nyu.edu