

Approximating generalized likelihood ratio tests with calibrated discriminative classifiers

Kyle Cranmer

KYLE.CRANMER@NYU.EDU

Center for Cosmology and Particle Physics

Center for Data Science

New York University

New York, NY 10003, USA

Editor:

Abstract

We demonstrate how discriminative classifiers can be used to approximate generalized likelihood ratio tests over high-dimensional data when a generative model for the data is available for training and calibration.

1. Introduction

In many areas of science, (generalized) likelihood ratio tests are established tools for statistical inference. Directly constructing the likelihood ratio for high-dimensional observations is often not possible or is computationally impractical. Here we demonstrate how discriminative classifiers can be used to construct equivalent likelihood ratio tests when a generative model for the data is available for training and calibration.

As a concrete example, consider searches for new particles at the Large Hadron Collider. The ATLAS and CMS experiments have published several hundred papers where the final statistical result was formulated as a hypothesis test or confidence interval using a generalized likelihood ratio test (Cowan et al., 2010). This includes the discovery of the Higgs boson (The ATLAS Collaboration, 2012; The CMS Collaboration, 2012) and subsequent measurement of its properties. The bulk of the likelihood ratio tests at the LHC are based on the distribution of a single event-level feature that discriminates between a hypothesized process of interest (labeled *signal*) and various other processes (labeled *background*).

To improve the statistical power of these tests, hundreds of these searches have utilized supervised learning to train discriminative classifiers. Within High Energy Physics (HEP) libraries such as TMVA that implement conventional techniques like the multi-layer perceptron and boosted decision trees (Hocker et al., 2007) are commonly used. Recently there has been progress in using deep networks (Baldi et al., 2014) and a NIPS workshop synthesizing the lessons learned during HiggsML (?), the largest Kaggle challenge in history.

As noted in (Whiteson and Whiteson, 2007), “such methods are suboptimal because they assume that the selector with the highest classification accuracy will yield a mass measurement with the smallest statistical uncertainty.” A key point is that evaluating the loss for classification is a per-event operation, while evaluating the loss for the mass measurement is a per-experiment operation involving on a dataset with many events. They went on to demonstrate a computationally intensive stochastic optimization technique based on the per-experiment loss out performed the two stage selection-estimation process.

Generalized likelihood ratio tests enjoy a number of optimal properties in the asymptotic limit and are widely used even away from that limit. The form of the generalized likelihood ratio is explicitly a per-experiment object, thus naively it has the same computational challenges as other approaches that optimize a per-experiment loss. This paper shows how one can reformulate generalized likelihood ratio test in terms of a per-event classification stage and a calibration process. Moreover, one can utilize a generative statistical model for the data both to train a discriminative classifier with a supervised learning algorithm and for the calibration process.

1.1 Notation and Assumptions

We use the following notation:

- x : a vector of features for an event
- D : a dataset of $D = \{x_1, \dots, x_n\}$, where x_e are assumed to be i.i.d.
- θ : parameters of a statistical model
- $p(x|\theta)$: probability density (statistical model) for x given θ
- $s(x; \theta_0, \theta_1)$: real-valued discriminative classification score, parametrized by θ_0 and θ_1
- $p(s_{\theta_0, \theta_1}|\theta)$: The probability density for $s(x; \theta_0, \theta_1)$ implied by $p(x|\theta)$

We will assume the x_e are i.i.d., so that $p(D|\theta) = \prod_{e=1}^n p(x_e|\theta)$.

1.2 Prelude

In the setting where one is interested in simple hypothesis testing between a null $\theta = \theta_0$ against an alternate $\theta = \theta_1$, the Neyman-Pearson lemma states that the likelihood ratio

$$T(D) = \prod_{e=1}^n \frac{p(x_e|\theta_0)}{p(x_e|\theta_1)} \quad (1)$$

is the most powerful test statistic. In order to evaluate $T(D)$, one must be able to evaluate the probability density $p(x|\theta)$ at any value x . However, it is increasingly common in science that one has a complex simulation that can act as generative model for $p(x|\theta)$, but one

cannot evaluate the density directly. For instance, this is the case high energy physics where the simulation of particle detectors can only be done in the ‘forward mode’. This same setting has been considered by Clayton Scott and Xin Tong (2013).

The main result of this paper is that one can form an equivalent test based on

$$T'(D) = \prod_{e=1}^n \frac{p(s_e|\theta_0)}{p(s_e|\theta_1)} \quad (2)$$

if

$$s_e = s(x_e; \theta_0, \theta_1) = m(p(x_e|\theta_0)/p(x_e|\theta_1)) \quad (3)$$

where m is any strictly increasing or decreasing function. This result will be proven below. This allows us to recast the original likelihood ratio test into an alternate form in which supervised learning is used to train the discriminative classifier $s(x; \theta_0, \theta_1)$. The discriminative classifier can be trained with data $(x, y = 0)$ generated from $p(x|\theta_0)$ and $(x, y = 1)$ generated from $p(x|\theta_1)$. In Section 3 we extend this result to generalized likelihood ratio tests, where it will be useful to have the classifier parametrized in terms of (θ_0, θ_1) .

While the original goal for frequentist hypothesis testing is to make a decision to accept or reject the null hypothesis based on the entire dataset D , we are able to reformulate it such that the machine learning problem is a per-event classification problem. This follows from the fact that we assume the x_e to be i.i.d.

1.3 Comments on classification and frequentist hypothesis tests

Vast literature exists around generative and discriminative classifiers (Andrew Y. Ng). Typically, generative classifiers learn a model for the joint probability $p(x, y)$, of the inputs x and the classification label y , and predict $p(y|x)$ via Bayes rule. In contrast, discriminative classifiers model the posterior $p(y|x)$ directly. For classification tasks, one then thresholds on $p(y|x)$. In both cases this description in terms of a posterior requires a prior distribution for $p(y)$, which is either modeled explicitly or learned from the training data. This familiar formulation of classification may lead to some confusion in the setting of the current work.

The first possible source of confusion we wish to avoid is that here $p(x|\theta)$ is a generative *statistical model* for the features x , not a generative classifier. We think of the $p(x|\theta)$ along the lines of a traditional scientific theory, able to make predictions about x and being motivated by domain-specific considerations. For example, in the context of HEP $p(x|\theta)$ is based on quantum field theory, a detailed simulation of the particle detector, and data processing algorithms that transform raw sensor data into the feature vector x . Moreover, we are not attempting to learn the generative model $p(x|\theta)$, we are taking it as given and trying to learn the corresponding (generalized) likelihood ratio test.

The second possible source of confusion is that we are not directly interested in calibrating the classification score in terms of a per-event posterior probability $p(y|x)$. Instead, we are interested in the approximation of the per-experiment likelihood ratio, which might be used to calculate p-values defined by $P(T(D) > k|\theta)$.

Lastly, in the setting of frequentist hypothesis tests, we do not have a prior $\pi(\theta)$. While we can use the generative models to produce training data $(x, y = 0)$ generated from $p(x|\theta_0)$ and $(x, y = 1)$ generated from $p(x|\theta_1)$, the relative mix $p(y)$ is arbitrary. When $p(y = 0) = p(y = 1) = 1/2$, then

$$p(y = 1|x) = \frac{p(x|y = 1)}{p(x|y = 0) + p(x|y = 1)} = \frac{p(x|\theta_1)}{p(x|\theta_0) + p(x|\theta_1)}, \quad (4)$$

which is monotonic with the desired likelihood ratio $p(x|\theta_1)/p(x|\theta_0)$. Since the prior $p(y)$ is not needed for the target likelihood ratio test and because the classifier score $p(y|x)$ may not be well calibrated, we choose to denote the classifier score $s(x)$ and simply think of it as a deterministic dimensionality reduction map $s : X \rightarrow \mathbb{R}$. Similar points are made in (Clayton Scott).

2. Dimensionality reduction and calibration

The target hypothesis test is based on

$$\ln T = \sum_{e=1}^n \underbrace{\log \left[\frac{p(x_e|\theta_0)}{p(x_e|\theta_1)} \right]}_{q(x_e)}. \quad (5)$$

Here we see that the test statistic T for the experiment is composed of a sum over events of the per-event function $q(x)$. A sum over a monotonic, but non-linear function of $q(x)$ would not lead to an equivalent hypothesis test.

The important part of the per-event function $q(x)$ is that it defines iso-contours in the feature space x . As we will show, our goal is to learn a monotonic function of $p(x|\theta_0)/p(x|\theta_1)$, which will share the same iso-contours. Then the remaining challenge is to find the appropriate rescaling that gives back linear function of $q(x)$. Our claim is that the generative model $p(x|\theta)$ can be used to calibrate the density $p(s|\theta)$ and that

$$\ln T' = \sum_{e=1}^n \underbrace{\log \left[\frac{p(s_e|\theta_0)}{p(s_e|\theta_1)} \right]}_{q(s_e)}, \quad (6)$$

leads to an equivalent test.

For notational simplicity, let $p_0(x) = p(x|\theta_0)$, $p_1(x) = p(x|\theta_1)$, and $s(x) = s(x; \theta_1, \theta_0)$. The distribution of x totally determines the distribution of s . In the application at hand, the function s maps a high-dimensional feature vector x to \mathbb{R}^+ . Let Ω_c be the level set $\{x \mid s(x) = c\}$ and $\hat{n} = \nabla s(x)/|\nabla s(x)|$ be the orthonormal vector to Ω_c at the point x .

We need to show that for all x , the density

$$p(q_x|\theta) = \int dx \delta(q_x - q_x(x)) p(x|\theta) / |\hat{n} \cdot \nabla q_x| \quad (7)$$

is equal to the density

$$p(q_s|\theta) = \int dx \delta(q_s - q_s(s(x))) p(x|\theta) / |\hat{n} \cdot \nabla q_s|. \quad (8)$$

It is sufficient to show that $q_x(x) = q_s(s(x)) \forall x \in \Omega_c$. The function $q_s(s)$ is based on the induced densities $p_0(s)$ and $p_1(s)$. The induced density $p_1(c)$ is given by

$$p_1(c) = \int dx \delta(c - s(x)) p_1(x) = \int d\Omega_c p_1(x) / |\hat{n} \cdot \nabla s| \quad (9)$$

100 and a similar equation for $p_0(c)$.

101

Theorem 1: We have the following equality

$$\frac{p_1(c)}{p_0(c)} = \frac{p_1(x)}{p_0(x)} \quad \forall x \in \Omega_c. \quad (10)$$

Proof For $x \in \Omega_c$, we can factor out of the integral the constant $p_1(x)/p_0(x)$. Thus

$$p_1(c) = \int dx \delta(c - s(x)) p_1(x) = \int d\Omega_c p_1(x) / |\hat{n} \cdot \nabla s| = \frac{p_1(x)}{p_0(x)} \int d\Omega_c p_0(x) / |\hat{n} \cdot \nabla s|, \quad (11)$$

and the integrals cancel in the likelihood ratio

$$\frac{p_1(c)}{p_0(c)} = \frac{p_1(x)}{p_0(x)} \frac{\int d\Omega_c p_0(x) / |\hat{n} \cdot \nabla s|}{\int d\Omega_c p_0(x) / |\hat{n} \cdot \nabla s|} = \frac{p_1(x)}{p_0(x)} \quad \forall x \in \Omega_c. \quad (12)$$

102 One can think of the ratio $p_1(s)/p_0(s)$ as a way of calibrating the the discriminative
103 classifier and correcting for the monotonic transformation m of the desired likelihood ratio
104 as in Eq. 3.

105 3. Composite hypotheses and the generalized likelihood ratio

In the case of composite hypotheses $\theta \in \Theta_0$ against an alternative $\theta \in \Theta_0^C$, the generalized likelihood ratio¹ test is commonly used

$$\lambda(x) = \frac{\sup_{\theta \in \Theta_0} p(D|\theta)}{\sup_{\theta \in \Theta} p(D|\theta)}. \quad (13)$$

106 This generalized likelihood ratio can be used both for hypothesis tests in the presence of
107 nuisance parameters or to create confidence intervals with or without nuisance parame-
108 ters. Often, the parameter vector is broken into two components $\theta = (\mu, \nu)$, where the μ

1. Also known as the profile likelihood ratio.

109 components are considered parameters of interest while the ν components are considered
 110 nuisance parameters. In that case Θ_0 corresponds to all values of ν with μ fixed.

Denote the maximum likelihood estimator

$$\hat{\theta} = \arg \max_{\theta} p(D|\theta) \quad (14)$$

and the conditional maximum likelihood estimator

$$\hat{\hat{\theta}} = \arg \max_{\theta \in \Theta_0} p(D|\theta) . \quad (15)$$

It is not obvious that if we are working with the distributions $p(s|\theta)$ (for some particular $s(x; \theta_0, \theta_1)$ comparison) that we can find the same estimators. Fortunately, there is a construction based on $p(s|\theta)$ that works. The maximum likelihood estimate of Eq. 14 is the same as the value that maximizes the likelihood ratio with respect to $p(D|\theta_1)$ for some fixed value of θ_1 . This allows us to use Theorem 1 to reformulate the maximum likelihood estimate

$$\hat{\theta} = \arg \max_{\theta} \frac{p(D|\theta)}{p(D|\theta_1)} = \arg \max_{\theta} \sum \ln \frac{p(x_e|\theta)}{p(x_e|\theta_1)} = \arg \max_{\theta} \sum \ln \frac{p(s(x_e; \theta, \theta_1)|\theta)}{p(s(x_e; \theta, \theta_1)|\theta_1)} . \quad (16)$$

111 It is important that we include the denominator $p(s(x_e; \theta, \theta_1)|\theta_1)$ because this cancels
 112 Jacobian factors that vary with θ .

113 4. Learning the correct mapping and its distribution

114 Thus far we have shown that likelihood ratio tests based on $p(x|\theta_0)/p(x|\theta_1)$ with high
 115 dimensional features x can be reproduced via hypothesis tests based on the univariate
 116 densities $p(s|\theta)$ for the particular dimensionality reduction map $s(x|\theta_0, \theta_1)$. In order for
 117 this approach to be useful in a setting where it is not possible to evaluate the density
 118 $p(x|\theta)$, we need to be able to approximate both $s(x|\theta_0, \theta_1)$ and $p(s|\theta)$ based on samples
 119 $\{x_i\}$ drawn from the generative model $p(x|\theta)$.

120 Denote the approximate dimensionality reduction map $\hat{s}(x; \theta_0, \theta_1)$ and its distribution
 121 $\hat{g}(\hat{s}|\theta)$. In general we will be interested in the machine learning problem that approximates
 122 these distributions based on samples $\{x_i\}$ drawn from the generative model $p(x|\theta)$. The first
 123 step in this direction is to confirm that a discriminative classifier obtained from common
 124 supervised learning algorithms will yield a function that is one-to-one with $p(x|\theta_0)/p(x|\theta_1)$.

125 4.1 The standard discriminative classification setting

126 For fixed θ_0 and θ_1 we can generate large samples from each model and train a classifier. To
 127 be concrete, let us use $p(x|\theta_0)$ to generate training data $(x_i, y_i = 0)$ and $p(x|\theta_1)$ to generate
 128 training data $(x_i, y_i = 1)$. With balanced training data $(p(y = 1) = p(y = 0) = 1/2)$

a quadratic loss function will lead to classifiers that approximate the regression function $\hat{s}(x) \approx p(y|x) = p(x|\theta_1)/(p(x|\theta_0) + p(x|\theta_1))$, which is monotonic with the desired per-event likelihood ratio $q(x)$. Thus, standard supervised learning algorithms with various surrogate loss functions lead to discriminative classifiers that approximate a monotonic function of per-event likelihood ratio $q(x)$. Once the classifier is trained, we can use the generative model together with a univariate density estimation technique (e.g. histograms or kernel density estimation) to approximate $\hat{g}(\hat{s}|\theta)$.

In the limit of large samples from the generative model, we can approximate the original likelihood ratio test arbitrarily well. With finite training data for $\hat{s}(x)$ and samples to approximate $\hat{g}(\hat{s}|\theta)$ it will be necessary to be more specific about the what loss function we are interested in for approximating the likelihood ratio test. This will depend in general on the ultimate goal of the test. We know that in the case of composite hypothesis tests that there is in general no uniformly most powerful test, thus it is likely that a decision theoretic approach taking into account some weighting or utility over the space Θ is necessary. This is left as a subject for future work.

4.2 Training a parametrized, discriminative classifier

We are left with the practical question of how to train a family of discriminative classifiers parametrized by θ_0 and θ_1 , the parameters associated to the null and alternate hypotheses, respectively. While this could be done independently for all θ_0 and θ_1 , it is desirable and convenient to have a smooth evolution of the classification score as a function of the parameters. Thus, we anticipate a single learning stage based on training data with input $(x, \theta_0, \theta_1)_i$ and target y_i . Somewhat unusually, the unknown values of the parameters are taken as input to the classifier, and are latent variables whose values will be specified via the enveloping (generalized) likelihood ratio test. We denote the learned family of classifiers $\hat{s}(x; \theta_0, \theta_1)$, and anticipate the training based roughly on the following algorithmic flow.

Algorithm 1 Training of the parametrized classifier.

```

initialize trainingData = {}
for  $\theta_0$  in  $\Theta$  do
  for  $\theta_1$  in  $\Theta$  do
    generate  $x_i^0 \sim p(x|\theta_0)$ 
    append  $\{(x_i^0, \theta_0, \theta_1, y = 0)\}$  to trainingData
    generate  $x_i^1 \sim p(x|\theta_1)$ 
    append  $\{(x_i^1, \theta_0, \theta_1, y = 1)\}$  to trainingData
  end for
end for
use trainingData to learn  $\hat{s}(x; \theta_0, \theta_1)$ 

```

154 While the function $p(x|\theta_1)/(p(x|\theta_0) + p(x|\theta_1))$ will minimize the expected squared loss
 155 based on training data produced according to Algorithm 1, it is not clear how training data
 156 from $\theta'_0 \neq \theta_0$ and $\theta'_1 \neq \theta_1$ will influence a real world classifier with finite capacity. This is
 157 left as an area for future work.

158 4.3 Embedding the classifier into the likelihood

159 In most settings that make use of likelihood ratio tests, the likelihood is based directly on
 160 some approximation of density for the observed data via $\hat{f}(x|\theta)$. In general, approximating
 161 the density $\hat{f}(x|\theta)$ is difficult for high-dimensional data, which motivates the use of the
 162 dimensionality reduction map $\hat{s}(x)$ and likelihood ratio tests based on the univariate density
 163 $\hat{g}(\hat{s}|\theta)$. In the case of a fixed classifier $\hat{s}(x)$ it is possible to pre-compute $\hat{s}_e = \hat{s}(x_e)$ and
 164 never refer back to the original features x_e . In the parametrized setting it is not possible
 165 to pre-compute $\hat{s}(x_e; \theta_0, \theta_1)$ for all values of θ_0 and θ_1 ; however, we can embed the classifier
 166 into the likelihood function to carry out the composition $\hat{g} \circ \hat{s}$. A concrete realization of
 167 this has been performed for probability models implemented with the RooFit probabilistic
 168 programming language and classifiers implemented with `scikit-learn` and TMVA (Verkerke
 169 and Kirkby, 2003; Pedregosa et al., 2011; Hocker et al., 2007).

170 In both cases, constructing the density $\hat{p}(\hat{s}|\theta)$ requires running the generative model at
 171 θ . In the context of a likelihood fit this would mean that the optimization algorithm that
 172 is trying to maximize the likelihood with respect to θ needs access to the generative model
 173 $p(x|\theta)$. This can be impractical when the generative model is computationally expensive
 174 or has high-latency (for instance some human intervention is required to reconfigure the
 175 generative model). In our HEP examples we have chosen to interpolate the distribution
 176 between discrete values of θ in order to produce a continuous parametrization for $\hat{p}(\hat{s}|\theta)$.
 177 One can easily imagine a number of approaches to embedding the classifier and estimating
 178 the density $\hat{p}(\hat{s}|\theta)$ and the relative merits of those approaches will depend critically on the
 179 dimensionality of θ and the computational cost of the generative model. We leave a more
 180 general strategy for this overarching optimization problem as an area of future work.

181 5. Applications in HEP

In HEP we are often searching for some class of events, generically referred to as *signal*,
 in the presence of a separate class of *background* events. For each event we measure some
 quantities x that have corresponding distributions $p_b(x|\nu)$ for background and $p_s(x|\nu)$ for
 signal, where ν are nuisance parameters describing uncertainties in the underlying physics
 prediction or response of the measurement device. The total model is a mixture of the
 signal and background, and μ is the mixture coefficient associate to the signal component.

$$p(D|\mu, \nu) = \prod_{e=1}^n [\mu p_s(x_e|\nu) + (1 - \mu) p_b(x_e|\nu)] , \quad (17)$$

182 New particle searches at the LHC are typically framed as hypothesis test where the null
 183 corresponds to $\mu = 0$, and the generalized likelihood ratio (profiling over ν as in Eq. 15) is
 184 used as a test statistic (Cowan et al., 2010). This approach was used for the discovery of
 185 the Higgs boson (The ATLAS Collaboration, 2012; The CMS Collaboration, 2012).

186 5.1 Typical usage pattern

187 In this setting, large samples of synthetic data $\{x_i, y_i\}$ generated with some nominal values
 188 of the parameters ν_0 , where $y = 0$ corresponds to the background density $p_b(x|\nu_0)$ and
 189 $y = 1$ corresponds to signal density $p_s(x|\nu_0)$. Importantly, the $y = 1$ label corresponds to
 190 the signal only, and not to the alternate signal-plus-background hypothesis. The resulting
 191 classifier approximates the regression function $p_s(x|\nu_0)/(p_s(x|\nu_0) + p_b(x|\nu_0))$, which is one
 192 to one with the likelihood ratio of the null to the alternate $p(x|\mu = 0, \nu_0)/p(x|\mu, \nu_0)$ for all
 193 μ . Associating the $y = 1$ label to the signal component has advantages because it helps the
 194 classifier focus its capacity on the relevant regions in the feature space, particularly when
 195 the signal is a very small perturbation to the background.

196 Once the classifier is trained, large samples of synthetic data drawn from $p_s(x|\nu)$ and
 197 $p_b(x|\nu)$, then we estimate the distributions $\hat{p}_s(\hat{s}|\nu)$ and $\hat{p}_b(\hat{s}|\nu)$ continuously parametrized
 198 in ν . An example of the distributions of the distribution of \hat{s} for the signal and background
 199 events with $\nu = \nu_0$ is shown in Figure 1.




Figure 1: Left: an example of the distributions $p_b(\hat{s}|\nu)$ and $p_s(\hat{s}|\nu)$ when the classifier s is a boosted-decision tree (BDT). Right: the corresponding ROC curve (right) for this and other classifiers. (Figures taken from TMVA manual.)

These steps feed into a subsequent statistical test based on the observed data $D = (x_1, \dots, x_n)$. For each event, the classifier is evaluated and one performs inference on a parameter μ related to the presence of the signal contribution. In particular, one forms

the statistical model²

$$p(D | \mu, \nu) = \prod_{e=1}^n [\mu \hat{p}_s(\hat{s}(x_e) | \nu) + (1 - \mu) \hat{p}_b(\hat{s}(x_e) | \nu)] . \quad (18)$$

200 5.2 Comments on typical usage of machine learning in HEP

201 Nuisance parameters are an after thought in the typical usage of machine learning in HEP.
 202 In fact, most discussions would related to the training and optimizing the classifier only
 203 consider $p_b(x)$ and $p_s(x)$ with $\nu = \nu_0$ being implicit. However, as experimentalists we know
 204 that we must account for various forms of systematic uncertainty, parametrized by nuisance
 205 parameters ν . In practice, we take the classifier as fixed and then propagate uncertainty
 206 through the classifier as in Eq. 18. Building the distribution $p(\hat{s} | \nu)$ for values of ν other
 207 than the nominal ν_0 used to train the classifier can be thought of as a calibration necessary
 208 for classical statistical inference; however, this classifier is clearly not optimal for $\nu \neq \nu_0$.

209 5.3 A more powerful approach

The standard use of machine learning in HEP can be improved by training a parametrized, discriminative classifier corresponding to the generalized likelihood ratio test

$$\lambda(\mu) = \frac{p(D | \mu, \hat{\nu})}{p(D | \hat{\mu}, \hat{\nu})} , \quad (19)$$

210 following the approach outlined in Section 3.

There is an interesting distinction between this approach and the standard use in which the classifier is trained for a fixed ν_0 . In the standard use one trains a classifier for signal vs. background, which is equivalent (in an ideal setting) to training a classifier for null (background-only) vs. alternate (signal-plus-background) since the resulting regression functions are one-to-one with each other. In contrast, in the case of the generalized likelihood ratio test

$$\frac{p(x | 0, \hat{\nu})}{p(x | \hat{\mu}, \hat{\nu})} = \frac{p_b(x | \hat{\nu})}{\hat{\mu} p_s(x_e | \hat{\nu}) + (1 - \hat{\mu}) p_b(x_e | \hat{\nu})} , \quad (20)$$

211 the background components don't cancel and there is an additional term $p_b(x | \hat{\nu})/p_b(x | \hat{\nu})$.
 212 In practice, with classifiers of finite capacity, there will be some tradeoff between taking
 213 into account this additional term and the more challenging learning problem when μ is
 214 very small.

2. Sometimes there is an additional Poisson term when expected number of signal and background events is known, which is referred to as an extended likelihood.

5.4 Decomposing tests between mixture models into their components

In this section we generalize the capacity focusing technique of training classifiers to discriminate between components of a mixture model. First, we generalize Eq. 17 to a mixture model of several components

$$p(x|\theta) = \sum_c w_c(\theta) p_c(x|\theta) . \quad (21)$$

It is possible to re-write the target likelihood ratio between two mixture models in terms of pairwise classification problems.

$$\frac{p(x|\theta_0)}{p(x|\theta_1)} = \frac{\sum_c w_c(\theta_0) p_c(x|\theta_0)}{\sum_{c'} w_{c'}(\theta_1) p_{c'}(x|\theta_1)} \quad (22)$$

$$= \sum_c \left[\sum_{c'} \frac{w_{c'}(\theta_1)}{w_c(\theta_0)} \frac{p_{c'}(x|\theta_1)}{p_c(x|\theta_0)} \right]^{-1} \quad (23)$$

$$= \sum_c \left[\sum_{c'} \frac{w_{c'}(\theta_1)}{w_c(\theta_0)} \frac{p_{c'}(s_{c,c',\theta_0,\theta_1}|\theta_1)}{p_c(s_{c,c',\theta_0,\theta_1}|\theta_0)} \right]^{-1} \quad (24)$$

The second line is a trivial, but useful decomposition into pair-wise classification between $p_{c'}(x|\theta_1)$ and $p_c(x|\theta_0)$. The third line uses Theorem 1 to relate the high-dimensional likelihood ratio into an equivalent calibrated likelihood ratio based on the univariate density of the corresponding classifier, denoted $s_{c,c',\theta_0,\theta_1}$. In the situation where the only free parameters of the mixture model are the coefficients w_c , then the distributions $p_c(s_{c,c',\theta_0,\theta_1}|\theta)$ are independent of θ and can be pre-computed (after training the discriminative classifier, but before performing the generalized likelihood ratio test). Equation 22 allows one to take advantage of both the parametrized classifier as in Eq. 20 and the capacity focusing technique in the typical HEP usage pattern.

5.5 Measuring particle properties

While the original motivation for this work was to improve the treatment of systematic uncertainties in new particle searches by parametrizing the classifier in terms of the nuisance parameters ν , the same approach can be used to for parameters of interest. In the case of new particle searches the parameter of interest is the mixture coefficient for the signal component $p_s(x|\nu)$. In the case of measuring particle properties the distribution of the features also depend on parameters of interest, such as a particle's mass and quantum numbers, so that one has $p_s(x|\nu) \rightarrow p_s(x|\theta)$, where θ includes both parameters of interest and nuisance parameters.

This is a significant step forward in the usage of machine learning in HEP, because classifiers have always been used between two static classes of events and not parametrized explicitly in terms of the physical quantities we wish to measure. The work of (Whiteson

239 and Whiteson, 2007) is the closest example as the stochastic optimization was directly
 240 trying to minimize the measurement uncertainty of a particle’s mass; however, the resulting
 241 classifier was fixed. This approach also offers the advantage that it explicitly relates the
 242 per-experiment optimization to the per-event optimization, which is less computationally
 243 intensive.

244 5.6 Related work

245 Clayton Scott and Xin Tong (2013) consider the machine learning problem associated to
 246 Neyman-Pearson hypothesis testing. As in this work, they consider the situation where
 247 one does not have access to the underlying distributions, but only has i.i.d. samples from
 248 each hypothesis. This work generalizes that goal from the Neyman-Pearson setting to gen-
 249 eralized likelihood ratio tests and emphasizes the connection with classification. Perhaps
 250 a formal treatment similar to the Neyman-Pearson case can be brought to bear in this
 251 more general setting. Tommi Jaakkola explore a way of leveraging generative models to
 252 derive kernel functions for use in discriminative methods. This interesting work is distinct
 253 from the point made here in which the generative model is being used for the purpose
 254 of providing training data and calibration. Rajat Raina and McCallum (2003) consider a
 255 hybrid generative/discriminative classifier; however, the goal of that work is not to leverage
 256 a generative model for the data, but to use both approaches to learn different subsets of
 257 the parameters in a single hybrid classifier. Bianca Zadrozny emphasize the importance of
 258 calibrated probability estimates from decision trees and naive Bayesian classifiers and in-
 259 vestigate various approaches to achieve this. In contrast to that work, we are not interested
 260 in calibrated probability estimates for $p(y|x)$ for individual events, but instead we use the
 261 calibration to correct for non-linear transformations of the target likelihood ratio and, per-
 262 haps, to provide calibrated p-values based on those likelihood ratio tests. Ihler et al. (2004)
 263 take on a different problem (tests of statistical independence) by using machine learning
 264 algorithms to find scalar maps from the high-dimensional feature space that achieve the
 265 desired statistical goal when the fundamental high-dimensional test is intractable.

266 6. Conclusions

267 We have shown that a parametrized family of discriminative classifiers $s(x; \theta_0, \theta_1)$ trained
 268 and calibrated with a generative model $p(x|\theta)$ can be used to approximate statistical in-
 269 ference likelihoods based on the ratio $p(x|\theta_0)/p(x|\theta_1)$ when it is not possible to evaluate the
 270 densities $p(x|\theta)$ for an arbitrary x . This approach leverages the power of machine learning
 271 in a classical statistical setting.

Acknowledgements

KC would like to thank Daniel Whiteson for discussions and encouragement throughout the project, Alex Ihler for challenging discussions that led to a reformulation of the initial idea, and Shimon Whiteson for patient feedback in that process. KC would also like to thank Yann LeCun, Philip Stark, and Pierre Baldi for their feedback on the project early in its conception, Balázs Kégl for discussions about the Kaggle challenge and feedback to the draft, and Yuri Shirman for reassuring cross checks of the Theorem. KC is supported by the US National Science Foundation grants PHY-0854724 and PHY-0955626. KC is grateful to UC-Irvine for their hospitality while this research was carried out and the Moore and Sloan foundations for their generous support of the data science environment at NYU.

References

- Michael I. Jordan Andrew Y. Ng. On Discriminative vs. Generative classifiers: A comparison of logistic regression and naive Bayes. URL <http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.19.9829&rank=1>.
- Pierre Baldi, Peter Sadowski, and Daniel Whiteson. Searching for Exotic Particles in High-Energy Physics with Deep Learning. *Nature Commun.*, 5:4308, 2014. doi: 10.1038/ncomms5308.
- Charles Elkan Bianca Zadrozny. Obtaining calibrated probability estimates from decision trees and naive Bayesian classifiers. URL <http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.29.3039>.
- Robert Nowak Clayton Scott. A Neyman-Pearson approach to statistical learning. URL <http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.103.6850&rank=5>.
- Glen Cowan, Kyle Cranmer, Eilam Gross, and Ofer Vitells. Asymptotic formulae for likelihood-based tests of new physics. *Eur.Phys.J.*, C71:1554, July 2010. doi: 10.1140/epjc/s10052-011-1554-0. URL <http://arxiv.org/abs/1007.1727>.
- Andreas Hocker, J. Stelzer, F. Tegenfeldt, H. Voss, K. Voss, et al. TMVA - Toolkit for Multivariate Data Analysis. *PoS, ACAT*:040, 2007.
- A.T. Ihler, J.W. Fisher, and A.S. Willsky. Nonparametric Hypothesis Tests for Statistical Dependency. *IEEE Transactions on Signal Processing*, 52(8):2234–2249, August 2004. ISSN 1053-587X. doi: 10.1109/TSP.2004.830994. URL <http://ieeexplore.ieee.org/lpdocs/epic03/wrapper.htm?arnumber=1315943>.
- F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau,

305 M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python.
 306 *Journal of Machine Learning Research*, 12:2825–2830, 2011.

307 Andrew Y. Ng Rajat Raina, Yirong Shen and Andrew McCallum. Clas-
 308 sification with hybrid generative/discriminative models. 2003. URL
 309 http://works.bepress.com/andrew_mccallum/38.

310 The ATLAS Collaboration. Observation of a new particle in the search for the Standard
 311 Model Higgs boson with the ATLAS detector at the LHC. *Phys.Lett.*, B716:1–29, 2012.
 312 doi: 10.1016/j.physletb.2012.08.020.

313 The CMS Collaboration. Observation of a new boson at a mass of 125 GeV
 314 with the CMS experiment at the LHC. *Phys.Lett.*, B716:30–61, 2012. doi:
 315 10.1016/j.physletb.2012.08.021.

316 David Haussler Tommi Jaakkola. Exploiting Generative Models in Discriminative Classi-
 317 fiers. URL <http://citeseer.ist.psu.edu/viewdoc/summary?doi=10.1.1.44.7709>.

318 Wouter Verkerke and David P. Kirkby. The RooFit toolkit for data modeling. *eConf*,
 319 C0303241:MOLT007, 2003.

320 S. Whiteson and D. Whiteson. Stochastic optimization for collision selection in high energy
 321 physics. *Association for the Advancement of Artificial Intelligence*, 292, 2007.

322 Xin Tong. A Plug-in Approach to Neyman-Pearson Classification.
 323 *Journal of Machine Learning Research*, 14:3011–3040, 2013. URL
 324 <http://jmlr.org/papers/v14/tong13a.html>.